

Quantum Query Algorithms for the Constructive Diagonal Ramsey Theorem

Anonymous author

Anonymous affiliation

Abstract

The constructive diagonal Ramsey problem asks, given adjacency-oracle access to an N -vertex graph, for a clique or independent set of the order guaranteed by Ramsey's theorem. We give a bounded-error quantum algorithm that, for every $K \geq 2$ and $N \geq 4^{K-1}$, finds and verifies a homogeneous K -set using

$$O\left(2^K K \log \frac{K}{\eta}\right)$$

edge queries with failure probability at most η . At the Ramsey scale $N = 2^n$, this yields a homogeneous set of order $\lfloor n/2 \rfloor + 1$ using $O(\sqrt{N} \log N \log(\log N/\eta))$ queries, improving on the $O(N)$ queries of the explicit Erdős–Szekeres recursion and giving, to our knowledge, the first sublinear worst-case algorithm for the Ramsey relation, whose query complexity was raised as a question by Jain, Li, Robere, and Xun (FOCS 2024). The bound also shows that the $N^{1-o(1)}$ quantum lower bound reported there cannot hold at its stated parameters; we show that the cited ingredients compose to an $\Omega(N^{1/24})$ lower bound instead.

The algorithm runs the constructive recursion over implicit candidate sets. Each set is represented by a short conjunction of adjacency constraints and sampled using capped unknown-solution quantum search, and a scale-aware concentration schedule balances estimation accuracy against the increasing cost of sampling deeper sets. We complement the upper bound with an $\Omega(N^{1-1/\sqrt{2}})$ randomized lower bound, transported from the random-Painter analysis of online Ramsey numbers, which holds on the uniform distribution $G(N, 1/2)$. On that distribution a greedy quantum search uses only $\tilde{O}(N^{1/4})$ queries, giving a provable polynomial quantum speedup for Ramsey search on random graphs. We also give an estimation-free size-biased recursion and extend it to every fixed number of edge colours.

2012 ACM Subject Classification Theory of computation $\rightarrow$ Quantum query complexity

Keywords and phrases Diagonal Ramsey theorem, clique or independent set, quantum query complexity, adjacency oracle, implicit sets, total search

1 Introduction

Ramsey's theorem guarantees that every graph on $N \geq 4^{K-1}$ vertices contains a clique or an independent set of order K . Its classical constructive proof [8] repeatedly chooses a pivot and restricts the candidate set to one of the pivot's two colour neighbourhoods; truncating each retained neighbourhood to half its current size, it finds a homogeneous K -set with $O(N)$ adjacency queries. The corresponding total search problem, RAMSEY, has been studied in proof complexity and in TFNP [15, 16, 14, 20, 13]: the input

model	input	lower bound	upper bound
randomized	worst case	$\Omega(N^{1-1/\sqrt{2}})$ [Proposition 5.2]	$O(N)$ [8]
randomized	$G(N, 1/2)$	$\Omega(N^{1-1/\sqrt{2}})$ [5]	$\tilde{O}(N^{1/2})$ (greedy)
quantum	worst case	$\Omega(N^{1/24})$ [Section 6.3]	$\tilde{O}(N^{1/2})$ [Theorem 1.1]
quantum	$G(N, 1/2)$	open	$\tilde{O}(N^{1/4})$ [Corollary 5.4]

■ **Table 1** Bounded-error query complexity of finding a homogeneous set of order $\lfloor n/2 \rfloor + 1$ in a graph on $N = 2^n$ vertices, in the worst case and on the uniform distribution $G(N, 1/2)$; here $1 - 1/\sqrt{2} \approx 0.2929$. The randomized lower bound is the random-Painter estimate of [5] for online Ramsey numbers, previously cited only for deterministic algorithms; the quantum lower bound is what the collision reduction of [13] yields at multiplicity two. The greedy classical bound on $G(N, 1/2)$ is folklore.

graph on $N = 2^n$ vertices is given by a circuit or an oracle, and the target order is $K = n/2$. Jain, Li, Robere, and Xun [13, Section III] explicitly raise the question of the query complexity of this relation in the deterministic, randomized, and quantum models. This paper answers the quantum question up to a polynomial gap and clarifies the lower-bound record on both the quantum and the classical side.

The central observation is that the candidate sets of the constructive proof need not be materialized. After i rounds, membership in the current set is determined by the i adjacency constraints imposed by the previous pivots. These short predicates allow quantum search to sample from the same recursion at substantially smaller query cost. The threshold 4^{K-1} is the natural scale for this recursive construction.

1.1 Our results

We work in the promised valid-graph adjacency-oracle model. One coherent query returns the colour of a requested unordered vertex pair, and the output is a set of vertices inducing one colour. Our main result is the following.

► **Theorem 1.1** (Implicit-majority quantum Ramsey search). *Let $K \geq 2$, $N \geq 4^{K-1}$, and $0 < \eta < 1/2$. Given coherent edge-oracle access to a simple undirected graph on N vertices, there is a quantum algorithm which, with probability at least $1 - \eta$, outputs a K -clique or a K -vertex independent set using*

$$O\left(2^K K \log \frac{K}{\eta}\right)$$

edge queries in the worst case. The algorithm verifies every returned witness and may output $\perp$ on the exceptional event.

At the succinct Ramsey scale, the theorem gives the following form.

► **Corollary 1.2.** *Let $n \geq 2$, $N = 2^n$, and $0 < \eta < 1/2$. Given coherent edge-oracle access to a simple graph on N vertices, one can find a clique or independent set of order $\lfloor n/2 \rfloor + 1$ using*

$$O\left(\sqrt{N} \log N \log \frac{\log N}{\eta}\right)$$

66 *edge queries in the worst case.*

67 **Proof.** Set $K = \lfloor n/2 \rfloor + 1$. Then

$$68 \quad 4^{K-1} = 2^{2\lfloor n/2 \rfloor} \leq 2^n = N,$$

69 so Theorem 1.1 applies. Moreover, $2^K \leq 2\sqrt{N}$, $K = O(\log N)$, and $\log(K/\eta) =$
 70 $O(\log(\log N/\eta))$. Substituting these estimates into the bound of Theorem 1.1 proves
 71 the claim. $\blacktriangleleft$

72 To our knowledge, this is the first sublinear worst-case query algorithm for the
 73 Ramsey relation in any model. Table 1 places it among the known bounds at the
 74 Ramsey scale, including the further results of this paper, which we now describe.

75 **An estimation-free recursion.**

76 Sampling a uniform candidate and following the colour of its edge to the current pivot
 77 selects a colour class with probability proportional to its size. A one-dimensional
 78 survival recurrence shows that this size-biased recursion completes with probability
 79 greater than $1/2$ from 4^{K-1} vertices, which gives an $O(2^K K^2 \log K \log(1/\eta))$ -query
 80 algorithm without any majority estimation (Proposition 5.1). The same recurrence
 81 extends to every fixed number of edge colours (Corollary C.1).

82 **A randomized lower bound.**

83 Every randomized algorithm needs $\Omega(M^{1-1/\sqrt{2}})$ queries to find a homogeneous K -set
 84 in a graph on $M = 4^{K-1}$ vertices, and the bound holds on the uniform distribution
 85 $G(M, 1/2)$ (Proposition 5.2). This improves the $M^{1/4}$ black-box bound of [12, 14].
 86 The proof transports the random-Painter weight estimate of [5] to the query model;
 87 Appendix B gives a self-contained version of that estimate.

88 **A quantum speedup on random graphs.**

89 On $G(N, 1/2)$, a greedy quantum neighbourhood search finds a K -clique with $\tilde{O}(N^{1/4})$
 90 queries (Proposition 5.3). Together with the distributional form of the lower bound,
 91 this gives a polynomial separation between the quantum and randomized query
 92 complexities of the Ramsey relation on its natural hard distribution (Corollary 5.4).
 93 We are not aware of an earlier provable quantum speedup for this relation. In the
 94 worst case, no separation is claimed: the randomized complexity is only known to lie
 95 between $N^{1-1/\sqrt{2}}$ and N .

96 **The quantum lower-bound record.**

97 The paragraph of [13] that raises the query question also reports an $N^{1-o(1)}$ quantum
 98 lower bound at the default parameters, obtained by composing its collision reduction
 99 with the multicollision lower bound of [18]. Theorem 1.1 shows that no such bound
 100 can hold at these parameters. Section 6.3 aligns the two formulations and shows by

direct substitution that the cited ingredients compose to an exponent of at most $1/24$, with equality at multiplicity two; the resulting $\Omega(N^{1/24})$ is, to our knowledge, the best known quantum lower bound. The main theorems of [13] do not depend on the remark.

1.2 Technical overview

Both algorithms represent the candidate set after round i by the conjunction of the adjacency constraints generated by the first i pivots. A membership test therefore uses $O(i)$ edge queries. Capped unknown-solution search samples from such a set without first enumerating its vertices; conditional on success, permutation symmetry makes the sample exactly uniform. The precise sampler guarantee and its proof appear in Lemma 2.2 and Appendix A.

The first recursion removes majority estimation entirely. At each pivot, it samples a uniform candidate and follows the colour of the sampled edge. A colour class is therefore selected with probability proportional to its size. The survival probability obeys a one-dimensional recurrence, and Lemma 4.1 shows that the recursion completes with probability greater than $1/2$ from 4^{K-1} vertices. This yields the estimation-free bound in Proposition 5.1.

The sharper recursion estimates a majority colour at each pivot. Sampling a deep candidate set is exponentially more expensive than sampling an early one, so imposing the same estimation accuracy at every level wastes queries. We instead distribute a fixed total error budget across the recursion. Early majorities receive higher accuracy, while later majorities tolerate larger error. A dyadic schedule keeps the cumulative loss in candidate-set size bounded and balances the work across levels. Conditional concentration, the resulting density invariant, and exact output correctness are proved in Lemmas 4.2–4.4; the query summation in Section 5 then proves Theorem 1.1.

The lower bound and the random-graph speedup both rest on the fact that the common neighbourhood of a small vertex set in $G(N, 1/2)$ has predictable size. For the lower bound, a martingale weight argument in the style of [5] shows that an adaptive algorithm with T queries exposes a monochromatic K -set with probability at most $2^{-\binom{K}{2} + c(c-1)}(2T)^{K-c}$ for every $c \leq K/2$; optimizing c gives the exponent $2 - \sqrt{2}$. For the speedup, a union bound shows that every j -set has common neighbourhood of density about 2^{-j} , so greedy quantum search costs only $O(2^{j/2})$ predicate evaluations at depth j and the recursion has depth K rather than $2K - 3$.

Organization.

Section 2 defines the query model and the sampler, and Sections 3–5 present the algorithms, the proofs, the lower bound, and the random-graph speedup. Section 6 compares the result with prior work and with the reported lower bound, and Section 7 concludes with open questions. The appendices contain the sampler proof, the lower-bound and random-graph proofs, the multicolour proof, numerical illustrations, and the AI disclosure.

2 Oracle model and capped quantum sampling

Let $G: \binom{[N]}{2} \rightarrow \{0, 1\}$ be a promised simple undirected graph, extended symmetrically to ordered pairs with $G(u, u) = 0$. Its coherent edge oracle acts as

$$O_G|u, v, b\rangle = |u, v, b \oplus G(u, v)\rangle.$$

Query complexity counts calls to O_G . Reversible operations on vertex labels are free, as usual in the query model, and coherent access to O_G is part of the input specification.

The same asymptotic bound applies to the standard locally verifiable TFNP relation with arbitrary Boolean-circuit encodings.

► **Corollary 2.1** (Locally verifiable arbitrary encoding). *The standard Ramsey TFNP relation that accepts either a local invalidity certificate or a locally verified homogeneous set has the same $O(2^K K \log(K/\eta))$ quantum-query upper bound as Theorem 1.1.*

Proof. Given a circuit $C(u, v)$, define the canonical simple graph

$$A(u, v) = \begin{cases} C(\min\{u, v\}, \max\{u, v\}), & u \neq v, \\ 0, & u = v. \end{cases}$$

One query to A uses $O(1)$ queries to C , including uncomputation. Run Theorem 1.1 on A and let W be the returned homogeneous set. Query the diagonal and both ordered entries induced by W , using $O(K^2)$ additional queries. A self-loop or an asymmetric pair is a local invalidity certificate. Otherwise C agrees with A on W , so W is a valid clique or independent-set witness. The verification cost is absorbed by the bound in Theorem 1.1. ◀

Our algorithms use the following consequence of unknown-solution quantum search [2]. We include its proof in Appendix A because the analysis needs both bounded termination and exact conditional uniformity.

► **Lemma 2.2** (Capped uniform marked-item search). *Let $T \subseteq [M]$ have density at least $\lambda > 0$, and suppose its membership predicate costs q edge queries. There is a quantum block which uses at most $O(q/\sqrt{\lambda})$ edge queries, succeeds with probability at least an absolute constant $p_0 > 0$, and, conditional on verified success, returns an exactly uniform element of T . The same predetermined cap makes the block terminate if the density promise is false or T is empty.*

Moreover, m independent uniform samples can be collected, except with probability δ , using a predetermined $O(m + \log(1/\delta))$ blocks whenever the density promise holds.

The batch guarantee avoids a separate $\log(1/\delta)$ overhead for each sample. Every block starts from the uniform state, uses fresh internal randomness, measures a candidate, and verifies membership. Conditional on success, the returned vertex is uniform on T ; successful outputs from distinct blocks are independent.

For comparison, consider the explicit classical recursion on $M = 4^{K-1} = 2^{r+1}$ vertices, where $r = 2K - 3$. At level i , query all edges from a pivot to the other

6 Quantum Constructive Ramsey Search

177 $s_i = M/2^i$ candidates. A majority colour contains at least $s_i/2$ of the $s_i - 1$ non-pivot
 178 vertices; retain exactly $s_i/2$ of them. Then

$$179 \quad \sum_{i < r} (s_i - 1) < 2M,$$

180 and two vertices remain after r rounds. The quantum algorithms retain the entire
 181 selected colour class but access it only through its membership predicate.

3 Two implicit-set recursions

3.1 An estimation-free size-biased recursion

184 Set $M = 4^{K-1}$ and $r = 2K - 3$, restrict the input to its first M vertices, let $S_0 = [M]$,
 185 and fix any $v_0 \in S_0$. At step $i = 0, \dots, r - 1$, define $T_i = S_i \setminus \{v_i\}$. Apply the repeated
 186 capped search from Lemma 2.2, with the universal density lower bound $1/M$, to
 187 sample an exactly uniform $x_i \in T_i$; abort if the fixed cap is exhausted. Query

$$188 \quad c_i = G(v_i, x_i),$$

189 and set

$$190 \quad S_{i+1} = \{x \in T_i : G(v_i, x) = c_i\}, \quad v_{i+1} = x_i.$$

191 After the final step, put $w = v_r$. Among the r labels c_i , choose a colour that appears
 192 at least $K - 1$ times. Return the corresponding pivots together with w , after directly
 193 verifying the induced edges.

194 The algorithm never materializes S_i . Its membership predicate is the conjunction
 195 of the edge constraints imposed by the earlier pivots, together with the exclusions of
 196 those pivots, and costs $O(i + 1)$ edge queries. Repeating each capped block $O(\log(2r))$
 197 times reduces its conditional failure probability on a nonempty set to $O(1/r)$. The
 198 probability that all r implicit sets remain nonempty is analyzed in Lemma 4.1.

3.2 Scale-aware approximate majorities

200 The sharper recursion estimates a near-majority colour at each pivot. Fix $M = 4^{K-1}$
 201 input vertices and set

$$202 \quad r = 2K - 3, \quad E = \frac{1}{16}.$$

203 For $0 \leq i < r$, define

$$204 \quad d_i = \left\lceil \frac{r - 1 - i}{6} \right\rceil, \quad z_i = 2^{-d_i}, \quad Z = \sum_{j=0}^{r-1} z_j,$$

205 and allocate

$$206 \quad \varepsilon_i = E \frac{z_i}{Z}, \quad a_i = \frac{1}{2} - \varepsilon_i. \tag{1}$$

Thus $\sum_i \varepsilon_i = E$. The allowable error doubles once every six levels, matching the increasing cost of sampling the deeper candidate sets.

Before round i , the recorded pivots $v_0, \dots, v_{i-1}$ and colours $c_0, \dots, c_{i-1}$ define the exact implicit set

$$S_i = \{u \in [M] \setminus \{v_0, \dots, v_{i-1}\} : G(v_j, u) = c_j \text{ for every } j < i\}. \quad (2)$$

Its reversible membership predicate uses $O(i)$ edge queries.

At round i , use Lemma 2.2 to obtain a verified pivot $v_i \in S_i$. Then collect, with replacement,

$$m_i = \left\lceil \frac{1}{2\varepsilon_i^2} \log \frac{4r}{\eta} \right\rceil \quad (3)$$

independent uniform samples from $S_i \setminus \{v_i\}$. Query their edges to v_i , and let $\hat{p}_i$ be the observed fraction of colour 1. Set $c_i = 1$ when $\hat{p}_i \geq 1/2$ and $c_i = 0$ otherwise. Appending this edge constraint to Equation (2) defines S_{i+1} .

After r rounds, one final capped search returns $w \in S_r$. Select a colour that occurs at least $K - 1$ times among $c_0, \dots, c_{r-1}$, take any $K - 1$ pivots carrying that colour, and add w . Query all $\binom{K}{2}$ induced edges and return the set only if this direct check confirms homogeneity.

Every search and sampling batch has a predetermined cap. Hence every execution terminates even if an inaccurate earlier estimate invalidates a later density promise. A failed search or batch returns $\perp$.

4 Correctness and survival of the nested sets

We first analyze the estimation-free recursion. For integers $d, s \geq 0$, let $P_d(s)$ be the infimum, over all graphs, candidate sets of size s , and distinguished pivots in those sets, of the probability that the ideal size-biased recursion completes d additional steps. Set $P_0(s) = 1$ for $s \geq 1$, $P_d(0) = 0$, and $P_d(1) = 0$ for $d \geq 1$.

► **Lemma 4.1** (Size-biased survival). *For every $d \geq 0$ and $s \geq 1$,*

$$P_d(s) \geq \frac{(s - 2^d + 1)_+}{s}, \quad (y)_+ = \max\{y, 0\}.$$

In particular, for $M = 4^{K-1}$ and $r = 2K - 3$, the ideal recursion completes with probability at least $1/2 + 1/M$.

Proof. The case $d = 0$ is immediate. Suppose $d \geq 1$ and $s \geq 2$, and let the two colour classes among the $s - 1$ non-pivot vertices have sizes a and b . A uniform sample selects these classes with probabilities $a/(s - 1)$ and $b/(s - 1)$. The induction hypothesis therefore gives

$$P_d(s) \geq \frac{aP_{d-1}(a) + bP_{d-1}(b)}{s - 1} \geq \frac{(a - 2^{d-1} + 1)_+ + (b - 2^{d-1} + 1)_+}{s - 1}.$$

Because $(x)_+ + (y)_+ \geq (x + y)_+$ and $a + b = s - 1$, the numerator is at least $(s - 2^d + 1)_+$. Replacing the denominator $s - 1$ by s weakens the bound and proves the claim; the

242 definitions cover $s \leq 1$. Finally, $M = 2^{r+1}$, so the bound at $(d, s) = (r, M)$ is
 243 $(M - 2^r + 1)/M = 1/2 + 1/M$. ◀

244 Exact conditional uniformity of each successful capped search couples an implemen-
 245 ted run to the ideal recursion. Assign each of the r searches failure probability at most
 246 $1/(100r)$. A union bound decreases the completion probability by at most $1/100$, so a
 247 single fixed-cap run succeeds with probability at least 0.49 . Independent repetition,
 248 followed by exact verification, amplifies this probability to $1 - \eta$ using $O(\log(1/\eta))$
 249 runs. On every completed run, nesting gives $G(v_i, v_j) = G(v_i, w) = c_i$ for $j > i$; the
 250 pigeonhole argument in Lemma 4.4 then gives a homogeneous output.

251 We next analyze the scale-aware recursion. The first lemma controls all adaptive
 252 estimates with one event.

253 ► **Lemma 4.2** (Conditional near-majorities). *Except with probability at most $\eta/2$, every*
 254 *completed sampling batch reached after successful searches and an accurate prefix*
 255 *satisfies*

$$256 \quad |\hat{p}_i - p_i| \leq \varepsilon_i,$$

257 where p_i is the true colour-1 fraction from v_i to $S_i \setminus \{v_i\}$.

258 **Proof.** Condition on the complete classical history $\mathcal{H}_i$ before the samples at level
 259 i . The graph, S_i , and v_i are then fixed. Let $\mathcal{B}_i$ be the event that the verified batch
 260 succeeds. Conditional on $\mathcal{B}_i$, Lemma 2.2 gives independent uniform samples from the
 261 fixed set. Hoeffding's inequality [11] and Equation (3) imply

$$262 \quad \Pr(|\hat{p}_i - p_i| > \varepsilon_i \mid \mathcal{H}_i, \mathcal{B}_i) \leq 2e^{-2m_i\varepsilon_i^2} \leq \frac{\eta}{2r}.$$

263 Multiplying by $\Pr(\mathcal{B}_i \mid \mathcal{H}_i) \leq 1$ gives the same upper bound for a completed but
 264 inaccurate batch conditional on $\mathcal{H}_i$. Apply this bound to the first inaccurate completed
 265 level after an accurate prefix. The tower property removes the conditioning, and a
 266 union bound over the r possible levels completes the proof. The argument does not
 267 require independence across levels. ◀

268 There are r pivot searches, r sampling batches, and one final search. Give each
 269 operation failure probability at most $\eta/(4r + 2)$ using the capped amplification in
 270 Lemma 2.2. Applying these conditional bounds after good prefixes, as above, shows
 271 that their union has probability at most $\eta/2$. Together with Lemma 4.2, all searches
 272 and estimates used below are good with probability at least $1 - \eta$.

273 Let $s_i = |S_i|$ and

$$274 \quad A_0 = 1, \quad A_i = \prod_{j=0}^{i-1} a_j.$$

275 ► **Lemma 4.3** (Variable recurrence and density). *On every prefix satisfying Lemma 4.2,*

$$276 \quad s_i > MA_i - 1, \quad A_i \geq \frac{7}{8} 2^{-i}.$$

277 Consequently S_r is nonempty, and for every sampling level $i < r$,

$$278 \quad \frac{|S_i \setminus \{v_i\}|}{M} > \frac{3}{8} 2^{-i}.$$

279 The final set also satisfies $|S_r|/M \geq 2^{-(r+1)} \geq (3/8)2^{-r}$.

280 **Proof.** An accurate empirical majority retains at least the fraction $a_i = 1/2 - \varepsilon_i$ of
281 the non-pivot vertices. Hence

$$282 \quad s_{i+1} \geq a_i(s_i - 1). \tag{4}$$

283 Since $\sum_i \varepsilon_i = 1/16$ and $\prod_j (1 - x_j) \geq 1 - \sum_j x_j$ for $x_j \in [0, 1]$,

$$284 \quad A_i = 2^{-i} \prod_{j < i} (1 - 2\varepsilon_j) \geq 2^{-i} \left(1 - 2 \sum_{j < i} \varepsilon_j \right) \geq \frac{7}{8} 2^{-i}.$$

285 We prove $s_i > MA_i - 1$ by induction. The claim holds at $i = 0$. If it holds at level
286 i , then Equation (4) and $a_i < 1/2$ give

$$287 \quad s_{i+1} > a_i(MA_i - 2) = MA_{i+1} - 2a_i > MA_{i+1} - 1.$$

288 Because $M2^{-r} = 2$, we obtain $s_r > (7/8)2 - 1 = 3/4$, and therefore the integer s_r is
289 positive. For $i < r$, we have $M2^{-i} \geq 4$ and

$$290 \quad s_i - 1 > MA_i - 2 \geq \frac{7}{8}M2^{-i} - 2 \geq \frac{3}{8}M2^{-i}.$$

291 Finally, $s_r \geq 1$ and $1/M = 2^{-(r+1)}$, which proves the last assertion. $\blacktriangleleft$

292 **► Lemma 4.4** (Exact homogeneous output). *Whenever the algorithm reaches its final*
293 *verified output, that output is a K -clique or a K -vertex independent set.*

294 **Proof.** For $i < j$, the definition of S_{i+1} gives $G(v_i, v_j) = c_i$, and it also gives $G(v_i, w) =$
295 c_i . Among the $r = 2K - 3$ colours, one occurs at least $\lceil r/2 \rceil = K - 1$ times. The
296 corresponding pivots together with w therefore induce that colour. The final direct
297 check verifies the returned set independently of earlier search or estimation errors. $\blacktriangleleft$

298 **5 Query complexity, a classical lower bound, and random graphs**

299 **5.1 The estimation-free recursion**

300 **► Proposition 5.1** (Size-biased quantum recursion). *Under the hypotheses of The-*
301 *orem 1.1, the estimation-free algorithm returns a verified homogeneous K -set with*
302 *probability at least $1 - \eta$ using*

$$303 \quad O\left(K^2 2^K \log K \log \frac{1}{\eta}\right)$$

304 *edge queries in the worst case.*

Proof. A capped sample from a nonempty subset of the M -element universe costs $O(\sqrt{M})$ membership tests. Repeating the block $O(\log(2r))$ times reduces its failure probability to $O(1/r)$. Since the membership predicate at level i costs $O(i+1)$ edge queries, one fixed-cap run uses

$$O\left(\sqrt{M} \log(2r) \sum_{i=0}^{r-1} (i+1)\right) = O(K^2 2^K \log K)$$

edge queries. The survival analysis following Lemma 4.1 gives a constant success probability for each run, including capped-search failures and final verification. Repeating independent runs $O(\log(1/\eta))$ times proves the stated worst-case bound. $\blacktriangleleft$

5.2 The scale-aware recursion

Proof of Theorem 1.1. Set $M = 4^{K-1}$ and run the algorithm of Section 3 on any fixed set of M input vertices. The failure allocation in Section 4, together with Lemma 4.2, implies that all required searches succeed and all completed estimates are accurate with probability at least $1 - \eta$. On this event, Lemma 4.3 guarantees every density promise used by the sampler and ensures that S_r is nonempty. The final output is homogeneous by Lemma 4.4. Because the algorithm verifies the returned set directly, every non- $\perp$ output is valid even outside the good event.

It remains to bound the number of queries. By Lemma 4.3, a capped search block at level i uses $O(2^{i/2})$ membership tests. Computing and uncomputing membership and querying the sampled edge costs $O(i+1)$ edge queries. Define

$$b_i = (i+1)2^{i/2}.$$

The batch cap in Lemma 2.2 uses $O(m_i + \log(r/\eta)) = O(m_i)$ blocks. By Equation (3), the total cost of the sampling batches is therefore

$$O\left(\log \frac{r}{\eta} \sum_{i=0}^{r-1} b_i \varepsilon_i^{-2}\right). \tag{5}$$

At most six indices share any value of d_i , and hence

$$Z = \sum_i 2^{-d_i} < 6 \sum_{d \geq 0} 2^{-d} = 12.$$

Using Equation (1),

$$\sum_{i=0}^{r-1} b_i \varepsilon_i^{-2} = O\left(\sum_{i=0}^{r-1} (i+1) 2^{i/2} 4^{d_i}\right).$$

Write $h = r - 1 - i$. Since $4^{\lceil h/6 \rceil} \leq 4 \cdot 2^{h/3}$, the last display is at most

$$O\left(2^{r/2} \sum_{h=0}^{r-1} (r-h) 2^{-h/6}\right) = O(r 2^{r/2}).$$

334 The amplified pivot searches and the final search cost

$$335 \quad O\left(\log \frac{r}{\eta} \sum_{i=0}^r b_i\right) = O\left(r 2^{r/2} \log \frac{r}{\eta}\right),$$

336 and the $O(K^2)$ output verification is smaller. Substituting $r = 2K - 3$ into Equation (5)
337 gives

$$338 \quad O\left(2^K K \log \frac{K}{\eta}\right),$$

339 as claimed. ◀

340 Each search begins in the uniform state on the fixed power-of-two universe $[M]$.
341 The recorded pivot labels and colours define the reversible membership predicate for
342 S_i , so the algorithm requires no materialized representation of any candidate set.

343 5.3 A randomized classical lower bound

344 We complement the upper bound with a randomized classical lower bound. It is the
345 random-Painter estimate of Conlon, Fox, Grinshpun, and He [5, Theorem 1] for online
346 Ramsey numbers, transported to the query model. Because that estimate is a statement
347 about uniformly random colourings, the bound holds on the uniform distribution
348 $G(N, 1/2)$ and not only in the worst case. Appendix B gives a self-contained proof.

349 ► **Proposition 5.2** (Randomized classical lower bound). *Let $K \geq 8$ and let $T =$
350 $\lfloor 2^{(2-\sqrt{2})K-2} \rfloor$. For every $N \geq K$, every randomized algorithm that makes at most T
351 edge queries outputs a homogeneous K -set of $G \sim G(N, 1/2)$ with probability at most
352 $5/8$. Consequently, for $M = 4^{K-1}$, every randomized edge-query algorithm that finds
353 a homogeneous K -set with worst-case probability at least $2/3$ uses*

$$354 \quad \Omega\left(2^{(2-\sqrt{2})K}\right) = \Omega\left(M^{1-1/\sqrt{2}}\right)$$

355 queries.

356 The exponent $1 - 1/\sqrt{2} \approx 0.2929$ improves the $M^{1/4}$ black-box lower bound
357 of [12, 14]. Jain et al. [13, Section III] already point to [5] for the best known
358 deterministic lower bound; Proposition 5.2 records the randomized and distributional
359 form that the same analysis yields. It leaves the randomized worst-case complexity
360 between $\Omega(M^{1-1/\sqrt{2}})$ and the $O(M)$ of the explicit recursion, and it does not separate
361 Theorem 1.1 from every classical algorithm in the worst case.

362 5.4 A quantum speedup on random graphs

363 On the distribution $G(M, 1/2)$ of Proposition 5.2, a greedy quantum neighbourhood
364 search is much cheaper than the worst-case algorithm of Theorem 1.1: the common
365 neighbourhood of every small clique in a random graph has essentially its expected
366 size, so no majority has to be estimated and the recursion has depth K rather than
367 $2K - 3$.

► **Proposition 5.3** (Quantum greedy search on random graphs). *Let $K \geq 2$, $M = 4^{K-1}$, and $0 < \eta < 1/2$, and let $G \sim G(M, 1/2)$. There is a quantum algorithm which makes*

$$O\left(K2^{K/2} \log \frac{K}{\eta}\right)$$

edge queries on every input, never outputs a set that is not a K -clique, and outputs a K -clique with probability at least $1 - \eta - \pi_K$ over G and its internal randomness, where $\pi_K = 2M^{K-1}e^{-2^{K-2}/12}$ satisfies $\pi_K < e^{-100}$ for all $K \geq 14$.

The proof, given in Appendix B.3, shows that with probability at least $1 - \pi_K$ every j -set with $j < K$ has at least $M2^{-j-1}$ common neighbours, and then runs the capped search of Lemma 2.2 with density parameter 2^{-j-1} at level j . Combined with the distributional form of Proposition 5.2, this gives a provable quantum speedup for Ramsey search on random graphs.

► **Corollary 5.4** (Separation on random graphs at the Ramsey scale). *Let $n \geq 14$, $N = 2^n$, $K = \lfloor n/2 \rfloor + 1$, $0 < \eta < 1/2$, and $G \sim G(N, 1/2)$.*

1. *A quantum algorithm making $O(N^{1/4} \log N \log(\log N/\eta))$ edge queries outputs a homogeneous K -set of G with probability at least $1 - \eta - \pi_K$.*
2. *Every randomized algorithm making at most $\lfloor 2^{(2-\sqrt{2})K-2} \rfloor = \Omega(N^{1-1/\sqrt{2}})$ edge queries outputs a homogeneous K -set of G with probability at most $5/8$.*

Since $1/4 < 1 - 1/\sqrt{2}$, the bounded-error quantum query complexity of the Ramsey relation on $G(N, 1/2)$ is polynomially smaller than its bounded-error randomized query complexity.

Proof. The subgraph induced on the first $M = 4^{K-1} \leq N$ vertices is distributed as $G(M, 1/2)$, so the first part follows from Proposition 5.3, using $2^{K/2} \leq 2^{n/4+1/2}$ and $K = O(\log N)$. The second part is the distributional statement of Proposition 5.2, which applies because $K \geq 8$. For the final sentence, take $\eta = 1/4$ and $K \geq 14$, so that the quantum algorithm succeeds with probability greater than $2/3$ while $5/8 < 2/3$. ◀

The size-biased recursion also extends to every fixed number of edge colours. The statement and proof appear in Appendix C.1.

6 Related work and parameter comparison

6.1 Quantum algorithms for other Ramsey tasks

Several quantum approaches to Ramsey problems address a different input and output task from the one studied here. Work based on adiabatic computation, quantum annealing, or quantum counting searches for colourings that establish or test Ramsey-number bounds [9, 1, 24, 22, 23, 21]. In those formulations, the search space consists of labelled graph colourings. Our input is one fixed graph given by an adjacency oracle, and the goal is to find a homogeneous set guaranteed to exist in that graph.

Generic quantum algorithms for clique, independent set, and fixed-subgraph containment operate in a closer oracle model [6, 4, 17, 26]. Quantum backtracking supplies

another general search primitive [19]. Our algorithms additionally exploit the nested majority structure of the constructive Ramsey proof: each candidate set is represented by the constraints accumulated along the recursion and sampled as an implicit set. We are not aware of an earlier implementation of the constructive Erdős–Szekeres recursion by this mechanism in the adjacency-query model.

6.2 Constructive Ramsey search

Ramsey search also appears in proof complexity and TFNP. It is connected to weak and structured pigeonhole principles [15, 16], its white-box hardness follows from collision-resistant hashing [14], and it belongs to the polynomial long-choice class PLC [20]. The closest matching query relation is studied in [13], which shows that it is not black-box reducible to PIGEON and, in the paragraph “Query complexity of RAMSEY” of its Section III, raises the question of its query complexity in the deterministic, randomized, and quantum models. On the classical side, the online Ramsey game (see [5] and the references therein), in which Builder queries edges and Painter answers adversarially, is exactly the deterministic version of this question with an unbounded vertex set; the best known lower bound for it is the random-Painter bound of [5] used in Proposition 5.2.

6.3 Comparison with a reported quantum lower bound

Model and parameter alignment.

Definition II.6 of [13] uses $N = 2^n$ vertices, and the convention following Definition II.7 sets the default target to $K = n/2$. The unnumbered “Query complexity of RAMSEY” paragraph in Section III (p. 415) reports an $N^{1-o(1)}$ quantum lower bound, based on its Lemma III.1 and the multicollision lower bound of [18].

Away from the diagonal, the graph produced by Lemma III.1 from a function h and a fixed graph A_0 is

$$A_h(u, v) = \mathbf{1}[h(u) = h(v)] \vee A_0(h(u), h(v)), \quad A_h(u, u) = 0.$$

It is a symmetric simple graph. A coherent query uses $O(1)$ queries to the value oracle for h , including uncomputation of the two hash values. Thus the reduction lies within our valid-graph promise. Moreover, Corollary 2.1 gives the same upper bound for the locally verifiable arbitrary-circuit relation, and Corollary 1.2 returns $\lfloor n/2 \rfloor + 1$ vertices, from which one may discard a vertex to meet the default target. The graph promise, encoding, and output size therefore align in the two formulations.

Exponent obtained by direct substitution.

Let H denote the multicollision range size and G the number of vertices in the resulting Ramsey instance. For fixed multiplicity t , the sufficient condition in Theorem I.3 of [13] is $G \geq H^{4t}/4^t$. At the boundary $G = \Theta_t(H^{4t})$, the exponent in [18] is

$$\alpha_t = \frac{2^{t-1} - 1}{2^t - 1}.$$

443 The exponent transferred to G is therefore

$$444 \quad \frac{\alpha_t}{4t} = \frac{2^{t-1} - 1}{4t(2^t - 1)} \leq \frac{1}{24}, \quad t \geq 2.$$

445 Equality holds at $t = 2$. For $t \geq 3$, the inequality $\alpha_t < 1/2$ gives $\alpha_t/(4t) <$
 446 $1/(8t) \leq 1/24$. Increasing G relative to H only reduces the transferred exponent. The
 447 multicollision theorem is stated for fixed t , so the cited result also does not support a
 448 growing-multiplicity limit.

449 Direct composition of the two cited statements therefore yields an exponent at most
 450 $1/24$, rather than $1 - o(1)$. Theorem 1.1 rules out an $N^{1-o(1)}$ lower bound at these
 451 parameters: the default relation has quantum query complexity $O(\sqrt{N} \log N \log \log N)$.
 452 The remark is not used elsewhere in [13], and the black-box separation theorems of
 453 that paper are unaffected.

454 The lower bound that the reduction does give.

455 The case $t = 2$ of the composition is valid and yields, to our knowledge, the best known
 456 quantum lower bound. Let A_0 be a graph on H vertices with no homogeneous set
 457 of order $2 \log_2 H$, which exists by Erdős's bound $R(k, k) > 2^{k/2}$ [7]. For $G = H^8/16$,
 458 Lemma III.1 of [13] with $t = 2$ maps a function $h: [G] \rightarrow [H]$ to the graph A_h above on
 459 G vertices, and every homogeneous set of order $\frac{1}{2} \log_2 G = 4 \log_2 H - 2$ in A_h contains
 460 two vertices with the same h -value: an independent set of A_h has distinct h -values
 461 and maps to an independent set of A_0 , so it has fewer than $2 \log_2 H$ vertices, while
 462 a clique of A_h maps to a clique of A_0 and therefore takes fewer than $2 \log_2 H$ values
 463 on $4 \log_2 H - 2$ vertices. Finding a collision in a random function $[G] \rightarrow [H]$ requires
 464 $\Omega(H^{1/3})$ quantum queries [25, 18], and a query to A_h costs $O(1)$ queries to h . Hence
 465 the default Ramsey relation on N vertices has quantum query complexity $\Omega(N^{1/24})$.
 466 Together with Corollary 1.2, its quantum query complexity lies between $\Omega(N^{1/24})$ and
 467 $O(\sqrt{N} \log N \log \log N)$.

468 7 Discussion and open questions

469 The query improvement rests on two complementary ideas. First, the algorithm stores
 470 each candidate set as the conjunction of the adjacency constraints created by earlier
 471 pivots. This representation supports quantum sampling without enumerating the
 472 surviving vertices. Second, the scale-aware analysis allocates estimation error according
 473 to the cost of the recursion level. The resulting schedule keeps the cumulative loss in
 474 candidate-set size bounded while concentrating most of the sampling effort at the less
 475 expensive early levels. Table 1 summarizes the resulting landscape; we close with the
 476 questions it leaves open.

477 Randomized worst case.

478 The randomized query complexity of the Ramsey relation on $N = 2^n$ vertices lies
 479 between $\Omega(N^{1-1/\sqrt{2}})$ (Proposition 5.2) and the $O(N)$ of the explicit recursion. We do

not know a randomized algorithm with $o(N)$ worst-case queries: classical rejection sampling from the implicit set at depth i costs 2^i membership tests, so the deep levels of the recursion cost $\Theta(N)$ classically, whereas amplitude amplification makes them cost $\Theta(\sqrt{N})$. A randomized lower bound of $N^{1/2+\Omega(1)}$ would separate the worst-case quantum and randomized complexities.

Quantum lower bounds.

By Section 6.3 and Corollary 1.2, the quantum query complexity lies between $\Omega(N^{1/24})$ and $O(\sqrt{N} \log N \log \log N)$. Is $\sqrt{N}$ optimal? The final levels of the recursion search sets of constant size inside a universe of size $\Theta(N)$, which costs $\Theta(\sqrt{N})$ for unstructured search, but these sets are defined by $O(\log N)$ adjacency constraints, and the optimality of Grover search does not apply to them.

Random graphs.

Corollary 5.4 separates the quantum and randomized complexities on $G(N, 1/2)$, but no nontrivial quantum lower bound specific to that distribution is known, so the quantum exponent $1/4$ may not be optimal. On the classical side, Conlon et al. [5] conjecture that the exponent $2 - \sqrt{2}$ in the lower bound $2^{(2-\sqrt{2})K}$ is not tight for cliques of fixed order and that the truth is $2^{(2/3+o(1))K}$, which would correspond to $N^{1/3+o(1)}$ at the Ramsey scale; their branch-and-fill strategy may also admit a quantum version with a smaller exponent than $1/4$.

Extensions.

The implicit-set method may extend to off-diagonal and hypergraph Ramsey search. Quantum walks give sublinear-query algorithms for fixed sub-hypergraph containment [10]; a Ramsey algorithm would additionally need to exploit the totality and recursive structure of the problem. For multiple edge colours, Corollary C.1 gives a size-biased algorithm, while a scale-aware analogue may yield a sharper dependence on the recursion depth.

Scope of the evidence.

Appendix C.2 reports finite experiments illustrating the survival recurrence and the conditional uniformity of the sampler. They are sanity checks; the query bounds follow from the analytic proofs in Sections 2–5 and the appendices. The paper claims query bounds only: it makes no claim about gate complexity, hardware, or a worst-case separation from every randomized algorithm.

References

- 1 Zhengbing Bian, Fabian Chudak, William G. Macready, Geordie Rose, and Frank Gaitan. Experimental determination of Ramsey numbers. *Physical Review Letters*, 111:130505, 2013. URL: <https://arxiv.org/abs/1201.1842>, arXiv:1201.1842, doi:10.1103/PhysRevLett.111.130505.
- 2 Michel Boyer, Gilles Brassard, Peter Høyer, and Alain Tapp. Tight bounds on quantum searching. *Fortschritte der Physik*, 46(4–5):493–506, 1998. URL: <https://arxiv.org/abs/quant-ph/9605034>, arXiv:quant-ph/9605034.
- 3 Herman Chernoff. A measure of asymptotic efficiency for tests of a hypothesis based on the sum of observations. *The Annals of Mathematical Statistics*, 23(4):493–507, 1952. doi:10.1214/aoms/1177729330.
- 4 Andrew M. Childs and Jason M. Eisenberg. Quantum algorithms for subset finding. *Quantum Information and Computation*, 5(7):593–604, 2005. arXiv:quant-ph/0311038, doi:10.26421/QIC5.7–7.
- 5 David Conlon, Jacob Fox, Andrey Grinshpun, and Xiaoyu He. Online Ramsey numbers and the subgraph query problem. In *Building Bridges II: Mathematics of László Lovász*, volume 28 of *Bolyai Society Mathematical Studies*, pages 159–194. Springer, 2019. URL: <https://arxiv.org/abs/1806.09726>, arXiv:1806.09726, doi:10.1007/978-3-662-59204-5_4.
- 6 Sebastian Dörn. Quantum complexity bounds for independent set problems. *arXiv preprint*, 2005. URL: <https://arxiv.org/abs/quant-ph/0510084>, arXiv:quant-ph/0510084.
- 7 Paul Erdős. Some remarks on the theory of graphs. *Bulletin of the American Mathematical Society*, 53(4):292–294, 1947. doi:10.1090/S0002-9904-1947-08785-1.
- 8 Paul Erdős and George Szekeres. A combinatorial problem in geometry. *Compositio Mathematica*, 2:463–470, 1935. URL: <https://eudml.org/doc/88611>.
- 9 Frank Gaitan and Lane Clark. Ramsey numbers and adiabatic quantum computing. *Physical Review Letters*, 108:010501, 2012. URL: <https://arxiv.org/abs/1103.1345>, arXiv:1103.1345.
- 10 François Le Gall, Harumichi Nishimura, and Seiichiro Tani. Quantum algorithms for finding constant-sized sub-hypergraphs. *Theoretical Computer Science*, 609:569–582, 2016. arXiv:1310.4127, doi:10.1016/j.tcs.2015.10.006.
- 11 Wassily Hoeffding. Probability inequalities for sums of bounded random variables. *Journal of the American Statistical Association*, 58(301):13–30, 1963. doi:10.1080/01621459.1963.10500830.
- 12 Russell Impagliazzo and Moni Naor. Decision trees and downward closures. In *Proceedings of the Third Annual Structure in Complexity Theory Conference*, pages 29–38. IEEE, 1988. doi:10.1109/SCT.1988.5260.
- 13 Siddhartha Jain, Jiawei Li, Robert Robere, and Zhiyang Xun. On pigeonhole principles and Ramsey in TFNP. In *2024 IEEE 65th Annual Symposium on Foundations of Computer Science (FOCS)*, pages 406–428. IEEE, 2024. arXiv:2401.12604, doi:10.1109/FOCS61266.2024.00033.

- 555 **14** Ilan Komargodski, Moni Naor, and Eylon Yogev. White-box vs. black-box
556 complexity of search problems: Ramsey and graph property testing. *Journal of*
557 *the ACM*, 66(5):34:1–34:28, 2019. doi:10.1145/3341106.
- 558 **15** Jan Krajíček. On the weak pigeonhole principle. *Fundamenta Mathematicae*,
559 170(1–2):123–140, 2001. doi:10.4064/fm170-1-8.
- 560 **16** Jan Krajíček. Structured pigeonhole principle, search problems and hard tautolo-
561 gies. *The Journal of Symbolic Logic*, 70(2):619–630, 2005. doi:10.2178/jsl/11
562 20224731.
- 563 **17** Troy Lee, Frédéric Magniez, and Miklos Santha. A learning graph based quantum
564 query algorithm for finding constant-size subgraphs. *Chicago Journal of Theoret-*
565 *ical Computer Science*, 2012(10), 2012. URL: <https://arxiv.org/abs/1109.5>
566 135, arXiv:1109.5135.
- 567 **18** Qipeng Liu and Mark Zhandry. On finding quantum multi-collisions. In *Advances*
568 *in Cryptology—EUROCRYPT 2019, Part III*, volume 11478 of *Lecture Notes*
569 *in Computer Science*, pages 189–218. Springer, 2019. arXiv:1811.05385,
570 doi:10.1007/978-3-030-17659-4_7.
- 571 **19** Ashley Montanaro. Quantum walk speedup of backtracking algorithms. *Theory*
572 *of Computing*, 14(15):1–24, 2018. URL: <https://arxiv.org/abs/1509.02374>,
573 arXiv:1509.02374, doi:10.4086/toc.2018.v014a015.
- 574 **20** Amol Pasarkar, Christos Papadimitriou, and Mihalis Yannakakis. Extremal
575 combinatorics, iterated pigeonhole arguments and generalizations of PPP. In
576 *14th Innovations in Theoretical Computer Science Conference (ITCS 2023)*,
577 volume 251 of *Leibniz International Proceedings in Informatics (LIPIcs)*, pages
578 88:1–88:20. Schloss Dagstuhl–Leibniz-Zentrum für Informatik, 2023. doi:10.4
579 230/LIPIcs.ITCS.2023.88.
- 580 **21** Joel E. Pion and Susan M. Mniszewski. Toward computing bounds for Ramsey
581 numbers using quantum annealing. *Quantum Information Processing*, 24(7):221,
582 2025. URL: <https://arxiv.org/abs/2311.04405>, arXiv:2311.04405,
583 doi:10.1007/s11128-025-04839-x.
- 584 **22** Ri Qu, Zong shang Li, Juan Wang, Yan ru Bao, and Xiao chun Cao. Computing
585 hypergraph Ramsey numbers by using quantum circuit. *Quantum Information*
586 *Processing*, 12(7):2487–2496, 2013. arXiv:1210.3419, doi:10.1007/s11128-0
587 13-0541-9.
- 588 **23** Mani Ranjbar, William G. Macready, Lane Clark, and Frank Gaitan. Generalized
589 Ramsey numbers through adiabatic quantum optimization. *Quantum Information*
590 *Processing*, 15(9):3519–3542, 2016. arXiv:1606.01078, doi:10.1007/s11128-0
591 16-1363-3.
- 592 **24** Hefeng Wang. Determining Ramsey numbers on a quantum computer. *Physical*
593 *Review A*, 93(3):032301, 2016. arXiv:1510.01884, doi:10.1103/PhysRevA.93.
594 032301.
- 595 **25** Mark Zhandry. A note on the quantum collision and set equality problems.
596 *Quantum Information and Computation*, 15(7–8):557–567, 2015. arXiv:1312.1
597 027, doi:10.26421/QIC15.7-8-2.

- 598 **26** Yechao Zhu. Quantum query complexity of constant-sized subgraph containment.
599 *International Journal of Quantum Information*, 10(3):1250019, 2012. **arXiv:**
600 1109.4165, doi:10.1142/S0219749912500190.

A Proof of the capped-search lemma

Proof of Lemma 2.2. Begin in the uniform superposition on $[M]$ and apply the randomized-iteration unknown-solution search of [2]. When the marked density is at least λ , truncate the geometric iteration schedule once its maximum iteration count reaches a sufficiently large constant multiple of $1/\sqrt{\lambda}$. The cited analysis gives a success probability $p_0 > 0$, independent of M and λ , using $O(1/\sqrt{\lambda})$ calls to the membership predicate. Because the cap is fixed before any oracle outcome is observed, the block terminates for every predicate.

For each fixed number of Grover iterations, all marked basis states have the same amplitude. Randomizing the iteration count preserves this permutation symmetry. After measuring a candidate, evaluate the exact membership predicate and accept only a marked outcome. Conditional on acceptance, the output is therefore uniform on the marked set T . Independent blocks reinitialize all registers and random choices, so their accepted outputs are independent conditional on the block-success indicators.

To collect a batch, run

$$B = C(m + \log(1/\delta))$$

independent capped blocks, where the absolute constant C is chosen from p_0 . The number of accepted blocks stochastically dominates $\text{Bin}(B, p_0)$. A Chernoff lower-tail bound [3] makes the probability of fewer than m acceptances at most δ . Conditional on any success pattern with at least m acceptances, the first m accepted outputs are independent and uniform on T . ◀

B Proofs for the classical lower bound and the random-graph speedup

B.1 A weight bound for adaptive queries to a random graph

We record the random-Painter estimate of Conlon, Fox, Grinshpun, and He [5, proof of Theorem 1] in the form used in this paper. We follow their weight argument, taking the matching number of the queried graph as the potential; removing the critical edge from a maximum matching then leaves a matching of the smaller set, which keeps the recursion closed.

Throughout this subsection $V \geq 2$, and $G \sim G(V, 1/2)$ assigns independent uniform colours in $\{0, 1\}$ to the pairs of $[V]$. A deterministic algorithm queries T distinct pairs adaptively. Let $e_1, \dots, e_T$ be the queried pairs, let Q_t be the graph formed by $e_1, \dots, e_t$, and let $\mathcal{F}_t$ be the σ -algebra generated by the first t answers. Because the algorithm is deterministic, e_{t+1} is $\mathcal{F}_t$ -measurable.

For $U \subseteq [V]$ and $0 \leq t \leq T$, let $e_t(U)$ be the number of pairs inside U among $e_1, \dots, e_t$, and call U *consistent at time t* if all of these pairs have colour 1. Define

$$w(U, t) = \begin{cases} 2^{-\binom{|U|}{2} + e_t(U)}, & U \text{ is consistent at time } t, \\ 0, & \text{otherwise,} \end{cases}$$

the conditional probability, given $\mathcal{F}_t$, that U becomes a colour-1 clique when its remaining pairs are revealed. Let $\nu(U, t)$ be the matching number of $Q_t[U]$; it is

nondecreasing in t and increases by at most one when a pair is added. For integers $k \geq 2c \geq 0$ set

$$w_{k,c}(t) = \sum_{|U|=k, \nu(U,t) \geq c} w(U, t).$$

► **Lemma B.1** (Weight bound). *For all integers $k \geq 2c \geq 0$,*

$$\mathbb{E} w_{k,c}(T) \leq 2^{-\binom{k}{2} + c(c-1)} V^{k-2c} T^c.$$

Proof. *Martingale property.* Fix U . If $e_{t+1} \not\subseteq U$, then $w(U, t+1) = w(U, t)$. If $e_{t+1} \subseteq U$, then with probability $1/2$ the new pair has colour 1 and $w(U, t+1) = 2w(U, t)$, and otherwise $w(U, t+1) = 0$. Hence $\mathbb{E}[w(U, t+1) \mid \mathcal{F}_t] = w(U, t)$, and in particular $\mathbb{E} w(U, T) = w(U, 0) = 2^{-\binom{|U|}{2}}$. Summing over all k -sets gives the case $c = 0$:

$$\mathbb{E} w_{k,0}(T) = \binom{V}{k} 2^{-\binom{k}{2}} \leq V^k 2^{-\binom{k}{2}}.$$

Monotonicity. Since $\nu(U, \cdot)$ is nondecreasing, $w_{k,c}(t+1) \geq \sum_{\nu(U,t) \geq c} w(U, t+1)$, and each event $\{\nu(U, t) \geq c\}$ is $\mathcal{F}_t$ -measurable. Taking conditional expectations gives $\mathbb{E} w_{k,c}(t) \leq \mathbb{E} w_{k,c}(t+1)$.

Critical times. Let $c \geq 1$. For each U let $\tau(U) = \min\{t : \nu(U, t) \geq c\}$, a stopping time with $\tau(U) \geq 1$. Since $\{\tau(U) = t\} \in \mathcal{F}_t$ and $w(U, \cdot)$ is a martingale,

$$\mathbb{E} w_{k,c}(T) = \sum_{|U|=k} \sum_{t=1}^T \mathbb{E}[w(U, T) \mathbf{1}[\tau(U) = t]] = \sum_{t=1}^T \mathbb{E} w_{k,c}^*(t),$$

where

$$w_{k,c}^*(t) = \sum_{|U|=k, \tau(U)=t} w(U, t).$$

Charging. Fix t and write $e_t = \{x, y\}$. Suppose $|U| = k$ and $\tau(U) = t$. Then $\nu(U, t-1) = c-1$, $\nu(U, t) = c$, and $e_t \subseteq U$. Moreover, every maximum matching of $Q_t[U]$ contains e_t , since a maximum matching avoiding e_t would be a matching of size c in $Q_{t-1}[U]$. Fix such a matching M , let $M' = M \setminus \{e_t\}$, let A be the set of $2(c-1)$ vertices covered by M' , and let $U' = U \setminus \{x, y\}$. Then M' is a matching of $Q_{t-1}[U']$, so $\nu(U', t-1) \geq c-1$. We claim that at most $2(c-1)$ of the pairs between $\{x, y\}$ and U' belong to Q_{t-1} . Indeed, if $xa \in Q_{t-1}$ with $a \in U' \setminus A$, then $M' \cup \{xa\}$ is a matching of size c in $Q_{t-1}[U]$, a contradiction; the same holds for y . For each $ab \in M'$, the pairs xa and yb cannot both lie in Q_{t-1} , since $(M' \setminus \{ab\}) \cup \{xa, yb\}$ would again be a matching of size c in $Q_{t-1}[U]$; likewise xb and ya cannot both lie in Q_{t-1} . Hence at most two of xa, xb, ya, yb lie in Q_{t-1} , and the claim follows by summing over the $c-1$ edges of M' .

Consequently $e_t(U) \leq e_{t-1}(U') + 1 + 2(c-1)$, and U is consistent at time t only if U' is consistent at time $t-1$ and e_t has colour 1. Using $\binom{k}{2} - \binom{k-2}{2} = 2k-3$, we obtain

$$w(U, t) \leq 2^{-(2k-2c-2)} w(U', t-1) \quad \text{if } e_t \text{ has colour 1,}$$

and $w(U, t) = 0$ otherwise. The map $U \mapsto U'$ is injective because $U = U' \cup e_t$.
Therefore

$$w_{k,c}^*(t) \leq \mathbf{1}[e_t \text{ has colour } 1] 2^{-(2k-2c-2)} w_{k-2,c-1}(t-1).$$

The indicator is independent of $\mathcal{F}_{t-1}$ with mean $1/2$, while $w_{k-2,c-1}(t-1)$ is $\mathcal{F}_{t-1}$ -measurable. Taking expectations and using monotonicity,

$$\mathbb{E} w_{k,c}^*(t) \leq 2^{-(2k-2c-1)} \mathbb{E} w_{k-2,c-1}(T),$$

and summing over $t \leq T$ gives

$$\mathbb{E} w_{k,c}(T) \leq T 2^{-(2k-2c-1)} \mathbb{E} w_{k-2,c-1}(T).$$

Iteration. Applying the last inequality c times, the exponents $2(k-2i) - 2(c-i) - 1$ for $i = 0, \dots, c-1$ sum to $2kc - 3c^2$, and the case $c = 0$ gives

$$\mathbb{E} w_{k,c}(T) \leq T^c 2^{-(2kc-3c^2)} V^{k-2c} 2^{-\binom{k-2c}{2}} = 2^{-\binom{k}{2}+c(c-1)} V^{k-2c} T^c,$$

because $2kc - 3c^2 + \binom{k-2c}{2} = \binom{k}{2} - c(c-1)$. ◀

► **Corollary B.2** (Fully exposed monochromatic sets). *Let $K \geq 2$, let $V \geq K$, and let a deterministic algorithm make $T \geq 1$ adaptive edge queries to $G \sim G(V, 1/2)$. For every integer $0 \leq c \leq K/2$, the expected number of K -sets all of whose $\binom{K}{2}$ pairs have been queried and carry the same colour is at most*

$$2 \cdot 2^{-\binom{K}{2}+c(c-1)} (2T)^{K-c}.$$

Proof. Renaming vertices in order of first appearance in a query does not change the joint distribution of the queried graph and its colours, because the colours of distinct pairs are independent and uniform; the renamed algorithm queries pairs inside $[V']$ with $V' = \min\{V, 2T\}$, and the number of fully exposed monochromatic K -sets is unchanged. A fully exposed colour-1 K -set U satisfies $\nu(U, T) = \lfloor K/2 \rfloor \geq c$ and $w(U, T) = 1$, so the number of such sets is at most $w_{K,c}(T)$. Lemma B.1 with $V' \leq 2T$ bounds its expectation by $2^{-\binom{K}{2}+c(c-1)} (2T)^{K-2c} T^c \leq 2^{-\binom{K}{2}+c(c-1)} (2T)^{K-c}$. The same bound holds for colour 0 by symmetry. ◀

B.2 Proof of the randomized lower bound

Proof of Proposition 5.2. Let $\beta = 2 - \sqrt{2}$ and $\alpha = 1 - 1/\sqrt{2}$, so that $\beta = 2\alpha$ and $\alpha^2 - 2\alpha + 1/2 = 0$. Set

$$T = \lfloor 2^{\beta K-2} \rfloor, \quad c = \lfloor \alpha K \rfloor = \alpha K + \epsilon, \quad -1 < \epsilon \leq 0.$$

Since $\alpha < 1/2$, we have $0 \leq c \leq K/2$, and $K \geq 8$ gives $T \geq 1$.

Consider first a deterministic algorithm making at most T queries to $G \sim G(N, 1/2)$; padding with unused queries, we may assume that it makes exactly T . By Corollary B.2 and Markov's inequality, the probability that the queried pairs contain a fully exposed monochromatic K -set is at most

$$2 \cdot 2^{-\binom{K}{2}+c(c-1)} (2T)^{K-c}. \tag{6}$$

Since $2T \leq 2^{\beta K-1}$, the base-2 logarithm of Equation (6) is at most

$$1 - \binom{K}{2} + c(c-1) + (K-c)(\beta K-1) = 1 - \frac{K}{2} + \epsilon^2,$$

where the identity follows by substituting $c = \alpha K + \epsilon$ and $\beta = 2\alpha$ and using $\alpha^2 - 2\alpha + 1/2 = 0$. For $K \geq 8$ this quantity is at most -2 , so the event has probability at most $1/4$.

Now condition on the query transcript and suppose that the algorithm outputs a K -set whose internal pairs have not all been exposed in one colour. If the exposed internal pairs contain both colours, the output is not homogeneous. If they contain exactly one colour, at least one unexposed pair must match that colour, which has conditional probability at most $1/2$, because unexposed pairs are independent fair coins given the transcript. If no internal pair has been exposed, the success probability is $2^{1-\binom{K}{2}} \leq 1/2$ for $K \geq 3$. Thus the conditional success probability is at most $1/2$ outside the fully exposed event, and the success probability under $G(N, 1/2)$ is at most

$$\frac{1}{4} + \frac{3}{4} \cdot \frac{1}{2} = \frac{5}{8}.$$

A randomized algorithm making at most T queries is a distribution over such deterministic algorithms, so its success probability on $G \sim G(N, 1/2)$ is also at most $5/8$. Finally, an algorithm whose worst-case success probability on graphs with M vertices is at least $2/3 > 5/8$ must make more than T queries, and $M = 4^{K-1}$ gives

$$2^{\beta K} = \Theta(M^{\beta/2}) = \Theta(M^{1-1/\sqrt{2}}),$$

which proves the proposition. ◀

B.3 Proof of the random-graph speedup

Proof of Proposition 5.3. Let $M = 4^{K-1}$. For $W \subseteq [M]$ let

$$C(W) = \{u \in [M] \setminus W : G(w, u) = 1 \text{ for every } w \in W\}$$

be the common neighbourhood of W . Call G *expanding* if $|C(W)| \geq M2^{-j-1}$ for every $j \in \{1, \dots, K-1\}$ and every j -set W ; for $W = \emptyset$ the bound $|C(\emptyset)| = M$ holds trivially.

Random graphs are expanding. Fix $1 \leq j \leq K-1$ and a j -set W . Then $|C(W)| \sim \text{Bin}(M-j, 2^{-j})$ with mean $\mu = (M-j)2^{-j}$. Since $M = 4^{K-1} \geq 4(K-1) \geq 4j$, we have $\mu \geq \frac{3}{4}M2^{-j}$, hence $M2^{-j-1} \leq \frac{2}{3}\mu$ and $\mu \geq \frac{3}{4}M2^{-(K-1)} = \frac{3}{2} \cdot 2^{K-2}$. The Chernoff bound $\Pr[X \leq (1-\gamma)\mu] \leq e^{-\gamma^2\mu/2}$ [3] with $\gamma = 1/3$ gives

$$\Pr[|C(W)| < M2^{-j-1}] \leq e^{-\mu/18} \leq e^{-2^{K-2}/12}.$$

A union bound over the at most $\sum_{j=1}^{K-1} M^j \leq 2M^{K-1}$ sets W shows that G fails to be expanding with probability at most $\pi_K = 2M^{K-1}e^{-2^{K-2}/12}$.

742 *The algorithm.* Set $W_0 = \emptyset$. For $j = 0, \dots, K-1$, given the clique $W_j =$
 743 $\{v_1, \dots, v_j\}$, run the capped search of Lemma 2.2 on $[M]$ for the predicate “ $u \notin W_j$
 744 and $G(v_i, u) = 1$ for all $i \leq j$ ”, with density parameter $\lambda_j = 2^{-j-1}$, repeating the block
 745 $O(\log(K/\eta))$ times so that its failure probability is at most η/K whenever the density
 746 promise holds. The predicate costs $O(j+1)$ edge queries, including uncomputation.
 747 Let v_{j+1} be the verified output and $W_{j+1} = W_j \cup \{v_{j+1}\}$. If any search fails, output
 748 $\perp$; otherwise output W_K .

749 *Correctness and cost.* Every accepted vertex was verified to be adjacent to all earlier
 750 vertices, so a non- $\perp$ output is a K -clique for every input graph. If G is expanding, then
 751 the search at level j has marked density at least λ_j , and a union bound over the K
 752 searches shows that the algorithm outputs W_K with probability at least $1 - \eta$. Hence
 753 it succeeds with probability at least $1 - \eta - \pi_K$ over G and its internal randomness.
 754 Each capped block at level j uses $O(2^{j/2})$ predicate evaluations, so the total number
 755 of edge queries is

$$756 \quad O\left(\log \frac{K}{\eta} \sum_{j=0}^{K-1} (j+1)2^{j/2}\right) = O\left(K2^{K/2} \log \frac{K}{\eta}\right)$$

757 on every input, because every block has a predetermined cap. ◀

758 **C Multicolour extension, numerical illustrations, and AI disclosure**

759 **C.1 A multicolour extension of size-biased survival**

760 Suppose every edge has one of $q \geq 2$ colours, and define

$$761 \quad r = q(K-2) + 1, \quad B_0 = 0, \quad B_d = 1 + q + \dots + q^{d-1} = \frac{q^d - 1}{q - 1} \quad (d \geq 1).$$

762 Choose the power of two

$$763 \quad L_q = 2^{\lceil \log_2(2B_r) \rceil},$$

764 so $2B_r \leq L_q < 4B_r$.

765 ► **Corollary C.1** (Multicolour size-biased Ramsey search). *For $q, K \geq 2$, every q -*
 766 *edge-coloured complete graph on at least L_q vertices admits a bounded-error quantum*
 767 *algorithm that outputs a monochromatic K -clique using*

$$768 \quad O\left(r^2 \sqrt{L_q} \log(2r) \log \frac{1}{\eta}\right)$$

769 *colour-oracle queries in the worst case, where r , B_r , and L_q are as defined above.*

770 **Proof.** Run the size-biased recursion of Section 3 for r rounds on the first L_q vertices,
 771 now following the colour of a uniformly sampled edge among q possibilities. Let $P_d^{(q)}(s)$
 772 be the minimum probability that the ideal recursion completes d further rounds from
 773 a candidate set of size s and a fixed pivot. We claim that

$$774 \quad P_d^{(q)}(s) \geq \frac{(s - B_d)_+}{s}$$

for all $d \geq 0$ and $s \geq 1$. The case $d = 0$ is immediate. For $d \geq 1$ and $s \geq 2$, let the q colour classes among the $s - 1$ non-pivot vertices have sizes $a_1, \dots, a_q$. The induction hypothesis gives

$$P_d^{(q)}(s) \geq \frac{\sum_{j=1}^q a_j P_{d-1}^{(q)}(a_j)}{s-1} \geq \frac{\sum_{j=1}^q (a_j - B_{d-1})_+}{s-1}.$$

Since $\sum_j a_j = s - 1$ and $B_d = 1 + qB_{d-1}$, the numerator is at least $(s - B_d)_+$. This proves the claimed bound after weakening the denominator from $s - 1$ to s . The preceding display also gives the stronger bound $(s - B_d)_+/(s - 1)$ whenever $d \geq 1$ and $s \geq 2$. Consequently, starting from $L_q \geq 2B_r$, the ideal recursion survives all r rounds with probability strictly greater than $1/2$.

A membership predicate at level i contains i colour constraints and costs $O(i + 1)$ oracle queries. Using the universal density bound $1/L_q$, one capped sample costs $O((i + 1)\sqrt{L_q})$ colour queries. Repeating the capped block $O(\log(2r))$ times makes each sampling failure probability $O(1/r)$, so one fixed-cap run costs

$$O\left(\sqrt{L_q} \log(2r) \sum_{i=0}^{r-1} (i + 1)\right) = O\left(r^2 \sqrt{L_q} \log(2r)\right).$$

The union bound preserves a positive constant success probability. Since $r = q(K - 2) + 1$, one of the q recorded colours occurs at least $\lceil r/q \rceil = K - 1$ times. The corresponding pivots and the final vertex form a monochromatic K -clique by the same nesting argument as in Lemma 4.4. Independent repetition and direct verification amplify the success probability to $1 - \eta$ at an additional $O(\log(1/\eta))$ factor. ◀

C.2 Small-instance numerical illustrations

Ideal-sampler recursion. We ran the scale-aware recursion at $K = 3$, $N = 16$, and $r = 3$ for 1,000 trials in each of five graph families. The error schedule was $(\varepsilon_0, \varepsilon_1, \varepsilon_2) = (1/64, 1/64, 1/32)$, with sample counts $m_i = \lceil \log(600)/(2\varepsilon_i^2) \rceil$, corresponding to a total estimation-failure budget of 0.01. Pivots were sampled uniformly, majority-estimation samples were independent with replacement, and ties selected colour 1. We replaced quantum search by ideal uniform sampling to isolate the combinatorial recursion.

The deterministic inputs were the complete graph, the empty graph, and the balanced complete bipartite graph $K_{8,8}$. For the random inputs, one $G(16, 1/2)$ graph was held fixed across all 1,000 trials, while the final family drew a fresh $G(16, 1/2)$ graph in each trial. Every returned triple was verified; exhaustion or a nonhomogeneous output counted as failure.

Graph family on 16 vertices	failed/trials	minimum final $ S_3 $
complete	0/1000	13
empty	0/1000	13
balanced complete bipartite $K_{8,8}$	0/1000	6
one fixed $G(16, 1/2)$	0/1000	2
fresh $G(16, 1/2)$ per trial	0/1000	2

807 All 5,000 trials returned homogeneous triples. The final column records the smallest
808 observed value of $|S_3|$ in each family.

809 **Isolated Grover sampler.** We evolved a 16-dimensional state vector initialized in
810 the uniform superposition, using marked-set sizes 1, 2, 3, 4, 8, 15. The measurement
811 distribution was averaged over a uniformly chosen number of Grover iterations from zero
812 through four. The probability of measuring a marked item ranged from 0.4026746749
813 to 0.5973253251. Conditional on a marked outcome, the largest deviation from the
814 uniform distribution on marked items was zero to machine precision. This calculation
815 isolates the permutation symmetry used in Lemma 2.2.

816 **Finite split-tree survival.** Minimizing the ideal size-biased survival probability over
817 every admissible colour-class split gives exact finite-depth comparisons with the analytic
818 bounds. For nine binary levels starting from 1,024 candidates, the minimum was
819 approximately 0.5086135566, above the lower bound 513/1024. For four ternary levels
820 starting from 128 candidates, the minimum was approximately 0.7155732457, above
821 88/127.

822 C.3 AI Disclosure

823 We used ChatGPT and Claude as auxiliary tools for literature search, brainstorming,
824 proof checking, and building experiment scripts. The authors determined the research
825 direction and presentations, evaluated and revised all AI-assisted material, and take full
826 responsibility for the correctness, originality, and final content of the paper, including
827 every proof and reference.